

possibility that IBM had inserted a trapdoor in LUCIFER. In any case, the design decisions remained a mystery for many years.

In 1990, Eli Biham and Adi Shamir showed how their method of differential cryptanalysis could be used to attack DES. The DES algorithm involves 16 rounds; differential cryptanalysis would be more efficient than exhaustively searching all possible keys if the algorithm used at most 15 rounds. This indicated that perhaps the designers of DES had been aware of this type of attack. A few years later, IBM released some details of the design criteria, which showed that indeed they had constructed the system to be resistant to differential cryptanalysts. This cleared up at least some of the mystery surrounding the algorithm.

The DES has lasted for a long time, but is becoming outdated. Brute force searches (see Section 4.6), though expensive, can now break the system. Therefore, NIST replaced it with a new system in the year 2000. However, it is worth studying DES since it represents a popular class of algorithms and it has been one of the most frequently used cryptographic algorithms in history.

The DES is a block cipher; namely, it breaks the plaintext into blocks of 64 bits, and encrypts each block separately. The actual mechanics of how this is done is often called a Feistel system, after Horst Feistel, who was part of the IBM team that developed LUCIFER. In the next section, we give a simple algorithm that has many of the characteristics of this type of system, but is small enough to use as an example. In Section 4.3, we show how differential cryptanalysis can be used to attack this simple system. We give the DES algorithm in Section 4.4, and describe ways it is implemented in Section 4.5. Finally, in Section 4.6, we describe recent progress in breaking DES.

For an extensive discussion of block ciphers, see [Schneier].

4.2 A Simplified DES-Type Algorithm

The DES algorithm is rather unwieldy to use for examples, so in the present section we present an algorithm that has many of the same features, but is much smaller. Like DES, the present algorithm is a block cipher. Since the blocks are encrypted separately, we assume throughout the present discussion that the full message consists of only one block.

The message has 12 bits and is written in the form L_0R_0 , where L_0 consists of the first 6 bits and R_0 consists of the last 6 bits. The key K has 9 bits. The i th round of the algorithm transforms an input $L_{i-1}R_{i-1}$ to the output L_iR_i using an 8-bit key K_i derived from K .

The main part of the encryption process is a function $f(R_{i-1}, K_i)$ that




Figure 4.1: One Round of a Feistel System.

takes a 6-bit input R_{i-1} and an 8-bit input K_i and produces a 6-bit output. This will be described later.

The output for the i th round is defined as follows:

$$L_i = R_{i-1} \text{ and } R_i = L_{i-1} \oplus f(R_{i-1}, K_i),$$

where \oplus denotes XOR, namely bit-by-bit addition mod 2. This is depicted in Figure 4.1.

This operation is performed for a certain number of rounds, say n , and produces the ciphertext $L_n R_n$.

How do we decrypt? Start with $L_n R_n$ and switch left and right to obtain $R_n L_n$. (Note: This switch is built into the DES encryption algorithm, so it is not needed when decrypting DES.) Now use the same procedure as before, but with the keys K_i used in reverse order K_n, \dots, K_1 . Let's see how this works. The first step takes $R_n L_n$ and gives the output

$$[L_n] \quad [R_n \oplus f(L_n, K_n)].$$

We know from the encryption procedure that $L_n = R_{n-1}$ and $R_n = L_{n-1} \oplus f(R_{n-1}, K_n)$. Therefore,

$$\begin{aligned} [L_n] \quad [R_n \oplus f(L_n, K_n)] &= [R_{n-1}] \quad [L_{n-1} \oplus f(R_{n-1}, K_n) \oplus f(L_n, K_n)] \\ &= [R_{n-1}] \quad [L_{n-1}]. \end{aligned}$$

The last equality again uses $L_n = R_{n-1}$, so that $f(R_{n-1}, K_n) \oplus f(L_n, K_n)$ is 0. Similarly, the second step of decryption sends $R_{n-1} L_{n-1}$ to $R_{n-2} L_{n-2}$.

Continuing, we see that the decryption process leads us back to R_0L_0 . Switching the left and right halves, we obtain the original plaintext L_0R_0 , as desired.

Note that the decryption process is essentially the same as the encryption process. We simply need to switch left and right and use the keys K_i in reverse order. Therefore, both the sender and receiver use a common key and they can use identical machines (though the receiver needs to reverse left and right inputs).

So far, we have said nothing about the function f . In fact, any f would work in the above procedures. But some choices of f yield much better security than others. The type of f used in DES is similar to that which we describe next. It is built up from a few components.

The first function is an expander. It takes an input of 6 bits and outputs 8 bits. The one we use is given in Figure 4.2.




Figure 4.2: The Expander Function.

This means that the first input bit yields the first output bit, the third input bit yields both the fourth and the sixth output bits, etc. For example, 011001 is expanded to 01010101.

The main components are called S-boxes. We use two:

$$S_1 = \begin{bmatrix} 101 & 010 & 001 & 110 & 011 & 100 & 111 & 000 \\ 001 & 100 & 110 & 010 & 000 & 111 & 101 & 011 \end{bmatrix}$$

$$S_2 = \begin{bmatrix} 100 & 000 & 110 & 101 & 111 & 001 & 011 & 010 \\ 101 & 011 & 000 & 111 & 110 & 010 & 001 & 100 \end{bmatrix}.$$

The input for an S-box has 4 bits. The first bit specifies which row will be used: 0 for the first row, 1 for the second. The other 3 bits represent a binary number that specifies the column: 000 for the first column, 001 for the second, ..., 111 for the last column. The output for the S-box consists of the three bits in the specified location. For example, an input of 1010 for S_1 means we look at the second row, third column, which yields the output 110.

The key K consists of 9 bits. The key K_i for the i th round of encryption is obtained by using 8 bits of K , starting with the i th bit. For example, if

$K = 010011001$, then $K_4 = 01100101$ (after 5 bits, we reached the end of K , so the last 2 bits were obtained from the beginning of K).

We can now describe $f(R_{i-1}, K_i)$. The input R_{i-1} consists of 6 bits. The expander function is used to expand it to 8 bits. The result is XORed with K_i to produce another 8-bit number. The first 4 bits are sent to S_1 , and the last 4 bits are sent to S_2 . Each S-box outputs 3 bits, which are concatenated to form a 6-bit number. This is $f(R_{i-1}, K_i)$. We present this in Figure 4.3.

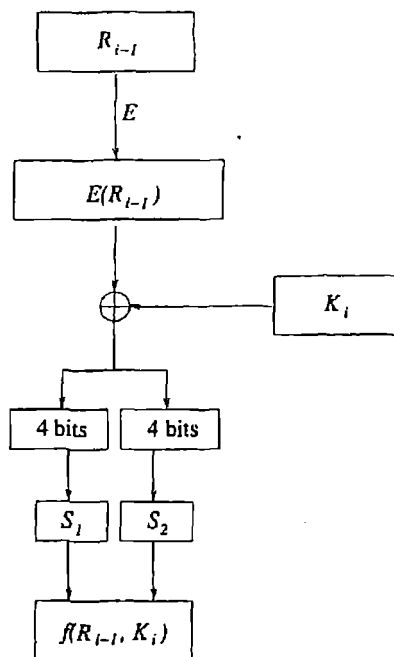


Figure 4.3: The Function $f(R_{i-1}, K_i)$.

For example, suppose $R_{i-1} = 100110$ and $K_i = 01100101$. We have

$$E(100110) \oplus K_i = 10101010 \oplus 01100101 = 11001111.$$

The first 4 bits are sent to S_1 and the last 4 bits are sent to S_2 . The second row, fifth column of S_1 contains 000. The second row, last column of S_2 contains 100. Putting these outputs one after the other yields $f(R_{i-1}, K_i) = 000100$.

We can now describe what happens in one round. Suppose the input is

$$L_{i-1}R_{i-1} = 011100100110$$

and $K_i = 01100101$, as previously. This means that $R_{i-1} = 100110$, as in the example just discussed. Therefore, $f(R_{i-1}, K_i) = 000100$. This is XORed with $L_{i-1} = 011100$ to yield $R_i = 011000$. Since $L_i = R_{i-1}$, we obtain

$$L_iR_i = 100110011000.$$

The output becomes the input for the next round.

4.3 Differential Cryptanalysis

This section is rather technical and can be skipped on a first reading.

Differential cryptanalysis was introduced by Biham and Shamir around 1990, though it was probably known much earlier to the designers of DES at IBM and NSA. The idea is to compare the differences in the ciphertexts for suitably chosen pairs of plaintexts and thereby deduce information about the key. Note that the difference of two strings of bits can be found by XORing them. Because the key is introduced by XORing with $E(R_{i-1})$, looking at the XOR of the inputs removes the effect of the key at this stage and hence removes some of the randomness introduced by the key. We'll see that this allows us to deduce information as to what the key could be.

4.3.1 Differential Cryptanalysis for Three Rounds

We eventually want to describe how to attack the above system when it uses four rounds, but we need to start by analyzing three rounds. Therefore, we temporarily start with L_1R_1 instead of L_0R_0 .

The situation is now as follows. We have obtained access to a three-round encryption device that uses the preceding procedure. We know all the inner workings of the encryption algorithm such as the S-boxes, but we do not know the key. We want to find the key by a chosen plaintext attack. We use various inputs L_1R_1 and obtain outputs L_4R_4 .

We have

$$R_2 = L_1 \oplus f(R_1, K_2)$$

$$L_3 = R_2 = L_1 \oplus f(R_1, K_2)$$

$$R_4 = L_3 \oplus f(R_3, K_4) = L_1 \oplus f(R_1, K_2) \oplus f(R_3, K_4).$$

Suppose we have another message $L_1^*R_1^*$ with $R_1 = R_1^*$. For each i , let $R_i' = R_i \oplus R_i^*$ and $L_i' = L_i \oplus L_i^*$. Then $L_i'R_i'$ is the "difference" (or sum; we are working mod 2) of L_iR_i and $L_i^*R_i^*$. The preceding calculation applied

to $L_1^* R_1^*$ yields a formula for R_4^* . Since we have assumed that $R_1 = R_1^*$, we have $f(R_1, K_2) = f(R_1^*, K_2)$. Therefore, $f(R_1, K_2) \oplus f(R_1^*, K_2) = 0$ and

$$R_4' = R_4 \oplus R_4^* = L_1' \oplus f(R_3, K_4) \oplus f(R_3^*, K_4).$$

This may be rearranged to

$$R_4' \oplus L_1' = f(R_3, K_4) \oplus f(R_3^*, K_4).$$

Finally, since $R_3 = L_4$ and $R_3^* = L_4^*$, we obtain

$$R_4' \oplus L_1' = f(L_4, K_4) \oplus f(L_4^*, K_4).$$

Note that if we know the input XOR, namely $L_1' R_1'$, and if we know the outputs $L_4 R_4$ and $L_4^* R_4^*$, then we know everything in this last equation except K_4 .

Now let's analyze the inputs to the S-boxes used to calculate $f(L_4, K_4)$ and $f(L_4^*, K_4)$. If we start with L_4 , we first expand and then XOR with K_4 to obtain $E(L_4) \oplus K_4$, which are the bits sent to S_1 and S_2 . Similarly, L_4^* yields $E(L_4^*) \oplus K_4$. The XOR of these is

$$E(L_4) \oplus E(L_4^*) = E(L_4 \oplus L_4^*) = E(L_1')$$

(the first equality follows easily from the bit-by-bit description of the expansion function). Therefore, we know

1. the XORs of the inputs to the two S-boxes (namely, the first four and the last four bits of $E(L_1')$);
2. the XORs of the two outputs (namely, the first three and the last three bits of $R_4' \oplus L_1'$).

Let's restrict our attention to S_1 . The analysis for S_2 will be similar. It is fairly fast to run through all pairs of 4-bit inputs with a given XOR (there are only 16 of them) and see which ones give a desired output XOR. These can be computed once for all and stored in a table.

For example, suppose we have input XOR equal to 1011 and we are looking for output XOR equal to 100. We can run through the input pairs (1011, 0000), (1010, 0001), (1001, 0010), ..., each of which has XOR equal to 1011, and look at the output XORs. We find that the pairs (1010, 0001) and (0001, 1010) both produce output XORs 100. For example, 1010 means we look at the second row, third column of S_1 , which is 110. Moreover, 0001 means we look at the first row, second column, which is 010. The output XOR is therefore $110 \oplus 010 = 100$.

We know L_4 and L_4^* . For example, suppose $L_4 = 101110$ and $L_4^* = 000010$. Therefore, $E(L_4) = 10111110$ and $E(L_4^*) = 00000010$, so the inputs

to S_1 are $1011 \oplus K_4^L$ and $0000 \oplus K_4^L$, where K_4^L denotes the left 4 bits of K_1 . If we know that the output XOR for S_1 is 100, then $(1011 \oplus K_4^L, 0000 \oplus K_4^L)$ must be one of the pairs on the list we just calculated, namely (1010, 0001) and (0001, 1010). This means that $K_4^L = 0001$ or 1010.

If we repeat this procedure a few more times, we should be able to eliminate one of the two choices for K_4 and hence determine 4 bits of K . Similarly, using S_2 , we find 4 more bits of K . We therefore know 8 of the 9 bits of K . The last bit can be found by trying both possibilities and seeing which one produces the same encryptions as the machine we are attacking.

Here is a summary of the procedure (for notational convenience, we describe it with both S-boxes used simultaneously, though in the examples we work with the S-boxes separately):

1. Look at the list of pairs with input XOR = $E(L'_4)$ and output XOR = $R'_4 \oplus L'_1$.
2. The pair $(E(L_4) \oplus K_4, E(L_4^*) \oplus K_4)$ is on this list.
3. Deduce the possibilities for K_4 .
4. Repeat until only one possibility for K_4 remains.

Example. We start with

$$L_1 R_1 = 000111011011$$

and the machine encrypts in three rounds using the key $K = 001001101$, though we do not yet know K . We obtain (note that since we are starting with $L_1 R_1$, we start with the shifted key $K_2 = 01001101$)

$$L_4 R_4 = 000011100101.$$

If we start with

$$L_1^* R_1^* = 101110011011$$

(note that $R_1 = R_1^*$), then

$$L_4^* R_4^* = 100100011000.$$

We have $E(L_4) = 00000011$ and $E(L_4^*) = 10101000$. The inputs to S_1 have XOR equal to 1010 and the inputs to S_2 have XOR equal to 1011. The S-boxes have output XOR $R'_4 \oplus L'_1 = 111101 \oplus 101001 = 010100$, so the output XOR from S_1 is 010 and that from S_2 is 100.

For the pairs (1001, 0011), (0011, 1001), S_1 produces output XOR equal to 010. Since the first member of one of these pairs should be the left four

bits of $E(L_4) \oplus K_4 = 0000 \oplus K_4$, the first four bits of K_4 are in $\{1001, 0011\}$. For the pairs $(1100, 0111), (0111, 1100)$, S_2 produces output XOR equal to 100. Since the first member of one of these pairs should be the right four bits of $E(L_4) \oplus K_4 = 0011 \oplus K_4$, the last four bits of K_4 are in $\{1111, 0100\}$.

Now repeat (with the same machine and the same key K) and with

$$L_1 R_1 = 010111011011 \text{ and } L_1^* R_1^* = 101110011011.$$

A similar analysis shows that the first four bits of K_4 are in $\{0011, 1000\}$ and the last four bits are in $\{0100, 1011\}$. Combining this with the previous information, we see that the first 4 bits of K_4 are 0011 and the last 4 bits are 0100. Therefore, $K = 00 * 001101$ (recall that K_4 starts with the fourth bit of K).

It remains to find the third bit of K . If we use $K = 000001101$, it encrypts $L_1 R_1$ to 001011101010, which is not $L_4 R_4$, while $K = 001001101$ yields the correct encryption. Therefore, the key is $K = 001001101$. ■

4.3.2 Differential Cryptanalysis for Four Rounds

Suppose now that we have obtained access to a four-round device. Again, we know all the inner workings of the algorithm except the key, and we want to determine the key. The analysis we used for three rounds still applies, but to extend it to four rounds we need to use more probabilistic techniques.

There is a weakness in the box S_1 . If we look at the 16 input pairs with XOR equal to 0011, we discover that 12 of them have output XOR equal to 011. Of course, we expect on the average that two pairs should yield a given output XOR, so the present case is rather extreme. A little variation is to be expected; we'll see that this large variation makes it easy to find the key.

There is a similar weakness in S_2 , though not quite as extreme. Among the 16 input pairs with XOR equal to 1100, there are 8 with output XOR equal to 010.

Suppose now that we start with randomly chosen R_0 and R_0^* such that $R_0' = R_0 \oplus R_0^* = 001100$. This is expanded to $E(001100) = 00111100$. Therefore the input XOR for S_1 is 0011 and the input XOR for S_2 is 1100. With probability $12/16$ the output XOR for S_1 will be 011, and with probability $8/16$ the output XOR for S_2 will be 010. If we assume the outputs of the two S-boxes are independent, we see that the combined output XOR will be 011010 with probability $(12/16)(8/16) = 3/8$. Because the expansion function sends bits 3 and 4 to both S_1 and S_2 , the two boxes cannot be assumed to have independent outputs, but $3/8$ should still be a reasonable estimate for what happens.

Now suppose we choose L_0 and L_0^* so that $L_0' = L_0 \oplus L_0^* = 011010$. Recall that in the encryption algorithm the output of the S-boxes is XORed

with L_0 to obtain R_1 . Suppose the output XOR of the S-boxes is 011010. Then $R'_1 = 011010 \oplus L'_0 = 000000$. Since $R'_1 = R_1 \oplus R'_1$, it follows that $R_1 = R'_1$.

Putting everything together, we see that if we start with two randomly chosen messages with XOR equal to $L'_0R'_0 = 011010001100$, then there is a probability of around $3/8$ that $L'_1R'_1 = 001100000000$.

Here's the strategy for finding the key. Try several randomly chosen pairs of inputs with XOR equal to 011010001100. Look at the outputs L_4R_4 and $L'_4R'_4$. Assume that $L'_1R'_1 = 001100000000$. Then use three-round differential cryptanalysis with $L'_1 = 001100$ and the known outputs to deduce a set of possible keys K_4 . When $L'_1R'_1 = 001100000000$, which should happen around $3/8$ of the time, this list of keys will contain K_4 , along with some other random keys. The remaining $5/8$ of the time, the list should contain random keys. Since there seems to be no reason that any incorrect key should appear frequently, the correct key K_4 will probably appear in the lists of keys more often than the other keys.

Here is an example. Suppose we are attacking a four-round device. We try one hundred random pairs of inputs L_0R_0 and $L'_0R'_0 = L_0R_0 \oplus 011010001100$. The frequencies of possible keys we obtain are in the following table. We find it easier to look at the first four bits and the last four bits of K_4 separately.

First 4 bits	Frequency	First 4 bits	Frequency
0000	12	1000	33
0001	7	1001	40
0010	8	1010	35
0011	15	1011	35
0100	4	1100	59
0101	3	1101	32
0110	4	1110	28
0111	6	1111	39
Last 4 bits	Frequency	Last 4 bits	Frequency
0000	14	1000	8
0001	6	1001	16
0010	42	1010	8
0011	10	1011	18
0100	27	1100	8
0101	10	1101	23
0110	8	1110	6
0111	11	1111	17

It is therefore likely that $K_4 = 11000010$. Therefore, the key K is $10^*110000$.

To determine the remaining bit, we proceed as before. We can compute that 000000000000 is encrypted to 100011001011 using $K = 101110000$ and is encrypted to 001011011010 using $K = 100110000$. If the machine we are attacking encrypts 000000000000 to 100011001011, we conclude that the second key cannot be correct, so the correct key is probably $K = 101110000$.

The preceding attack can be extended to more rounds by extensions of these methods. It might be noticed that we could have obtained the key at least as quickly by simply running through all possibilities for the key. That is certainly true in this simple model. However, in more elaborate systems such as DES, differential cryptanalytic techniques are much more efficient than exhaustive searching through all keys, at least until the number of rounds becomes fairly large. In particular, the reason that DES uses 16 rounds appears to be because differential cryptanalysis is more efficient than exhaustive search until 16 rounds are used.

There is another attack on DES, called **linear cryptanalysis**, that was developed by Mitsuru Matsui [Matsui]. The main ingredient is an approximation of DES by a linear function of the input bits. It is theoretically faster than an exhaustive search for the key and requires around 2^{43} plaintext-ciphertext pairs to find the key. It seems that the designers of DES had not anticipated linear cryptanalysis. For details of the method, see [Matsui].

4.4 DES

A block of ciphertext consists of 64 bits. The key has 56 bits, but is expressed as a 64-bit string. The 8th, 16th, 24th, ..., bits are parity bits, arranged so that each block of 8 bits has an odd number of 1s. This is for error detection purposes. The output of the encryption is a 64-bit ciphertext.

The DES algorithm, depicted in Figure 4.4, starts with a plaintext m of 64 bits, and consists of three stages:

1. The bits of m are permuted by a fixed initial permutation to obtain $m_0 = IP(m)$. Write $m_0 = L_0R_0$, where L_0 is the first 32 bits of m_0 and R_0 is the last 32 bits.
2. For $1 \leq i \leq 16$, perform the following:

$$\begin{aligned} L_i &= R_{i-1} \\ R_i &= L_{i-1} \oplus f(R_{i-1}, K_i), \end{aligned}$$

where K_i is a string of 48 bits obtained from the key K and f is a function to be described later.

3. Switch left and right to obtain $R_{16}L_{16}$, then apply the inverse of the initial permutation to get the ciphertext $c = IP^{-1}(R_{16}L_{16})$.

Decryption is performed by exactly the same procedure, except that the keys K_1, \dots, K_{16} are used in reverse order. The reason this works is the same as for the simplified system described in Section 4.2. Note that the left-right switch in step 3 of the DES algorithm means that we do not have to do the left-right switch that was needed for decryption in Section 4.2.

We now describe the steps in more detail.

The initial permutation, which seems to have no cryptographic significance, but which was perhaps designed to make the algorithm load more efficiently into chips that were available in 1970s, can be described by the Initial Permutation table. This means that the 58th bit of m becomes the 1st bit of m_0 , the 50th bit of m becomes the 2nd bit of m_0 , etc.

Initial Permutation															
58	50	42	34	26	18	10	2	60	52	44	36	28	20	12	4
62	54	46	38	30	22	14	6	64	56	48	40	32	24	16	8
57	49	41	33	25	17	9	1	59	51	43	35	27	19	11	3
61	53	45	37	29	21	13	5	63	55	47	39	31	23	15	7

The function $f(R, K_i)$, depicted in Figure 4.5, is described in several steps.

1. First, R is expanded to $E(R)$ by the following table.

Expansion Permutation												
32	1	2	3	4	5	4	5	6	7	8	9	
8	9	10	11	12	13	12	13	14	15	16	17	
16	17	18	19	20	21	20	21	22	23	24	25	
24	25	26	27	28	29	28	29	30	31	32	1	

This means that the first bit of $E(R)$ is the 32nd bit of R , etc. Note that $E(R)$ has 48 bits.

2. Compute $E(R) \oplus K_i$, which has 48 bits, and write it as $B_1 B_2 \dots B_8$, where each B_j has 6 bits.
3. There are 8 S-boxes S_1, \dots, S_8 , given on page 128. B_j is the input for S_j . Write $B_j = b_1 b_2 \dots b_6$. The row of the S-box is specified by $b_1 b_6$ while $b_2 b_3 b_4 b_5$ determines the column. For example, if $B_3 = 001001$, we look at the row 01, which is the second row (00 gives the first row) and column 0100, which is the 5th column (0100 represents 4 in binary; the first column is numbered 0, so the fifth is labeled 4). The entry in S_3 in this location is 3, which is 3 in binary. Therefore, the output of S_3 is 0011 in this case. In this way, we obtain eight 4-bit outputs C_1, C_2, \dots, C_8 .




Figure 4.4: The DES Algorithm.




Figure 4.5: The DES Function $f(R_{i-1}, K_i)$.

4. The string $C_1 C_2 \dots C_8$ is permuted according to the following table.

16	7	20	21	29	12	28	17	1	15	23	26	5	18	31	10
2	8	24	14	32	27	3	9	19	13	30	6	22	11	4	25

The resulting 32-bit string is $f(R, K_j)$.

Finally, we describe how to obtain K_1, \dots, K_{16} . Recall that we start with a 64-bit K .

1. The parity bits are discarded and the remaining bits are permuted by the following table.

Key Permutation															
57	49	41	33	25	17	9	1	58	50	42	34	26	18		
10	2	59	51	43	35	27	19	11	3	60	52	44	36		
63	55	47	39	31	23	15	7	62	54	46	38	30	22		
14	6	61	53	45	37	29	21	13	5	28	20	12	4		

Write the result as $C_0 D_0$, where C_0 and D_0 have 28 bits.

2. For $1 \leq i \leq 16$, let $C_i = LS_i(C_{i-1})$ and $D_i = LS_i(D_{i-1})$. Here LS_i means shift the input one or two places to the left, according to the following table.

Number of Key Bits Shifted per Round																
Round	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
Shift	1	1	2	2	2	2	2	2	1	2	2	2	2	2	2	1

3. 48 bits are chosen from the 56-bit string $C_i D_i$ according to the following table. The output is K_i .

14	17	11	24	1	5	3	28	15	6	21	10
23	19	12	4	26	8	16	7	27	20	13	2
41	52	31	37	47	55	30	40	51	45	33	48
44	49	39	56	34	53	46	42	50	36	29	32

It turns out that each bit of the key is used in approximately 14 of the 16 rounds.

A few remarks are in order. In a good cipher system, each bit of the ciphertext should depend on all bits of the plaintext. The expansion $E(R)$ is

S-Boxes

S-box 1

14	4	13	1	2	15	11	8	3	10	6	12	5	9	0	7
0	15	7	4	14	2	13	1	10	6	12	11	9	5	3	8
4	1	14	8	13	6	2	11	15	12	9	7	3	10	5	0
15	12	8	2	4	9	1	7	5	11	3	14	10	0	6	13

S-box 2

15	1	8	14	6	11	3	4	9	7	2	13	12	0	5	10
3	13	4	7	15	2	8	14	12	0	1	10	6	9	11	5
0	14	7	11	10	4	13	1	5	8	12	6	9	3	2	15
13	8	10	1	3	15	4	2	11	6	7	12	0	5	14	9

S-box 3

10	0	9	14	6	3	15	5	1	13	12	7	11	4	2	8
13	7	0	9	3	4	6	10	2	8	5	14	12	11	15	1
13	6	4	9	8	15	3	0	11	1	2	12	5	10	14	7
1	10	13	0	6	9	8	7	4	15	14	3	11	5	2	12

S-box 4

7	13	14	3	0	6	9	10	1	2	8	5	11	12	4	15
13	8	11	5	6	15	0	3	4	7	2	12	1	10	14	9
10	6	9	0	12	11	7	13	15	1	3	14	5	2	8	4
3	15	0	6	10	1	13	8	9	4	5	11	12	7	2	14

S-box 5

2	12	4	1	7	10	11	6	8	5	3	15	13	0	14	9
14	11	2	12	4	7	13	1	5	0	15	10	3	9	8	6
4	2	1	11	10	13	7	8	15	9	12	5	6	3	0	14
11	8	12	7	1	14	2	13	6	15	0	9	10	4	5	3

S-box 6

12	1	10	15	9	2	6	8	0	13	3	4	14	7	5	11
10	15	4	2	7	12	9	5	6	1	13	14	0	11	3	8
9	14	15	5	2	8	12	3	7	0	4	10	1	13	11	6
4	3	2	12	9	5	15	10	11	14	1	7	6	0	8	13

S-box 7

4	11	2	14	15	0	8	13	3	12	9	7	5	10	6	1
13	0	11	7	4	9	1	10	14	3	5	12	2	15	8	6
1	4	11	13	12	3	7	14	10	15	6	8	0	5	9	2
6	11	13	8	1	4	10	7	9	5	0	15	14	2	3	12

S-box 8

13	2	8	4	6	15	11	1	10	9	3	14	5	0	12	7
1	15	13	8	10	3	7	4	12	5	6	11	0	14	9	2
7	11	4	1	9	12	14	2	0	6	10	13	15	3	5	8
2	1	14	7	4	10	8	13	15	12	9	0	3	5	6	11

designed so that this will happen in only a few rounds. The purpose of the initial permutation is not completely clear. It has no cryptographic purpose. The S-boxes are the heart of the algorithm and provide the security. Their design was somewhat of a mystery until IBM published the following criteria in the early 1990s (for details, see [Coppersmith1]).

1. Each S-box has 6 input bits and 4 output bits. This was the largest that could be put on one chip in 1974.
2. The outputs of the S-boxes should not be close to being linear functions of the inputs (linearity would have made the system much easier to analyze).
3. Each row of an S-box contains all numbers from 0 to 15.
4. If two inputs to an S-box differ by 1 bit, the outputs must differ by 2 bits.
5. If two inputs to an S-box differ in their first 2 bits but have the same last 2 bits, the outputs must be unequal.
6. There are 32 pairs of inputs having a given XOR. For each of these pairs, compute the XOR of the outputs. No more than eight of these output XORs should be the same. This is clearly to avoid an attack via differential cryptanalysis.
7. A criterion similar to (6), but involving three S-boxes.

In the early 1970s, it took several months of searching for a computer to find appropriate S-boxes. Now, such a search could be completed in a very short time.

4.4.1 DES Is Not a Group

One possible way of effectively increasing the key size of DES is to double encrypt. Choose keys K_1 and K_2 and encrypt a plaintext P by $E_{K_2}(E_{K_1}(P))$. Does this increase the security?

Meet-in-the-middle attacks on cryptosystems are discussed in Section 4.7. It is pointed out that, if an attacker has sufficient memory, double encryption provides little extra protection. Moreover, if a cryptosystem is such that double encryption is equivalent to a single encryption, then there is no additional security obtained by double encryption.

In addition, if double encryption is equivalent to single encryption, then the (single encryption) cryptosystem is much less secure than one might guess initially (see Exercise 9 in Chapter 8). If this were true for DES, for